

# Brief **p4est** interface schematics

Carsten Burstedde

May 22, 2013

## Abstract

We describe the general procedure of using **p4est** from application codes. **p4est** is a software library that stores and modifies a forest-of-octrees refinement structure using distributed (MPI) parallelism. It expects the description of the domain as a coarse mesh of conforming hexahedra. Non-conforming adaptive mesh refinement (AMR), coarsening, and other operations that modify the forest are implemented as **p4est** API functions. To inform the application about the refinement structure, several methods are provided that encode this information.

## 1 Starting point

We generally separate the adaptive mesh refinement (AMR) topology from any numerical information. The former is stored and modified internally by the **p4est** software library, while an application is free to define the way it uses this information and arranges numerical and other data. This document is intended to describe the interface by which **p4est** transfers mesh information to the application.

The general, modular AMR pipeline is described in [3], which is not specific to **p4est** but can in principle be applied to any AMR provider. The **p4est** algorithms and main interface routines are described in [4]. An example usage of **p4est** as scalable mesh backend for the general-purpose finite element software **deal.II** is described in [1]. A reference implementation of **p4est** in **C** can be freely downloaded [2] and used and extended under the GNU General Public License. This software is best installed standalone into a dedicated directory, where an application code can then find the header and library files to compile and link against, respectively.

In this document, we document the three distinct tasks to

- A** create a coarse mesh (Figure 1),
- B** modify the refinement and partition structure internal to **p4est**,
- C** and to transfer the mesh information to an application.

Unless indicated otherwise, all operations described below are understood as MPI collectives, that is, they are called on all processors simultaneously. Currently, part A needs to be performed redundantly on all processors, which is acceptable for up to  $10^5$ – $10^6$  coarse cells (octree roots). In parts B and C, runtime and memory are roughly proportional to the number of elements (octree leaves) on a given processor, independent of the number of octrees.

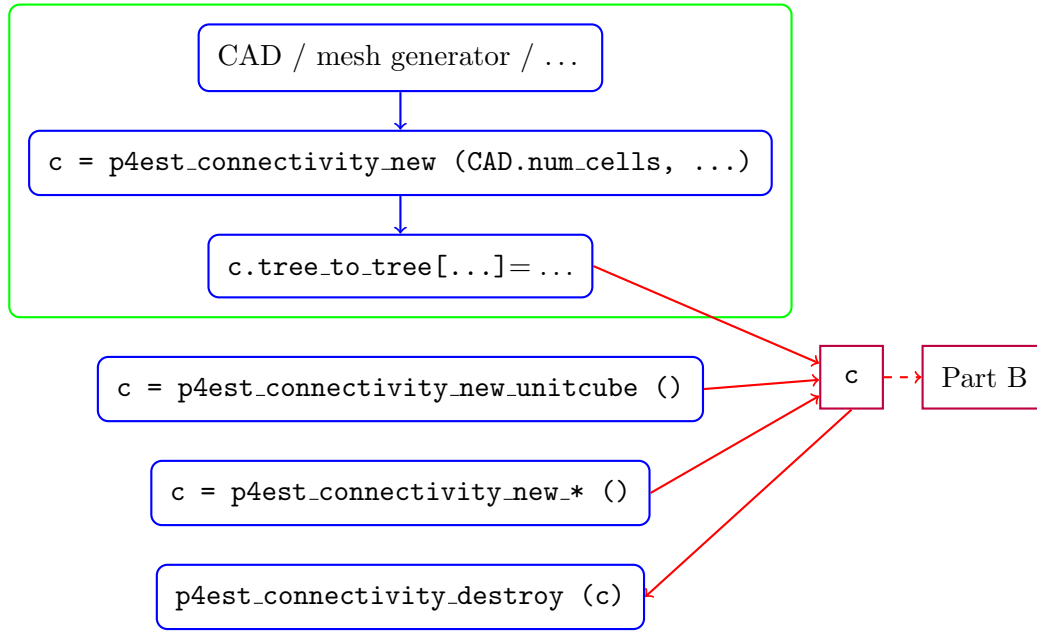


Figure 1: Part A, creating the coarse mesh connectivity. The connectivity `c` is a `C struct` that contains numbers and orientations of neighboring coarse cells. It can be created by translating CAD or mesh data file formats or by using one of several `p4est` convenience functions. The data format is documented in the big comment blocks in `p4est_connectivity.h` (2D) and `p8est_connectivity.h` (3D).

## References

- [1] W. BANGERTH, C. BURSTEDDE, T. HEISTER, AND M. KRONBICHLER, *Algorithms and data structures for massively parallel generic adaptive finite element codes*, ACM Transactions on Mathematical Software, 38 (2011), pp. 14:1–14:28.
- [2] C. BURSTEDDE, *p4est: Parallel AMR on forests of octrees*, 2010. <http://www.p4est.org/>.
- [3] C. BURSTEDDE, O. GHATTAS, G. STADLER, T. TU, AND L. C. WILCOX, *Towards adaptive mesh PDE simulations on petascale computers*, in Proceedings of Teragrid '08, 2008.
- [4] C. BURSTEDDE, L. C. WILCOX, AND O. GHATTAS, *p4est: Scalable algorithms for parallel adaptive mesh refinement on forests of octrees*, SIAM Journal on Scientific Computing, 33 (2011), pp. 1103–1133.